

Anshita Makode

Bangalore | anshitamakode@gmail.com | +91-9340072838 | <https://www.linkedin.com/in/anshita-makode/> | <https://github.com/anshitamakode>

PROFESSIONAL SUMMARY

AI/ML Engineer with 5 years of experience designing and shipping production multi-agent systems, LLM-powered features, and real-time data pipelines across retail, fintech, and EV domains. Deep hands-on expertise in agent orchestration, prompt optimization, and Generative AI on GCP/Vertex AI; pursuing an MS in Computer Science (AI/ML specialization) at UT Austin.

EDUCATION

The University of Texas at Austin

MS, Computer Science (AI/ML Specialization).

2027

Bharati Vidyapeeth's College of Engineering for Women, Pune

BE, Computer Engineering.

2020

EXPERIENCE

Software Engineer 2, AI/ML

May 2025 – Present

Best Buy

- Built and shipped agents for a production multi-agent product-discovery assistant (AskBlue) using Google ADK, the A2A protocol, and Gemini on Vertex AI, deployed as Cloud Run services on GCP; contributed across planner orchestration, NLU, and discovery/search.
- Implemented RAG-based grounding for the discovery flow using embedding-based semantic retrieval (Vertex AI Embeddings) over the product catalog, anchoring LLM responses to real inventory to prevent hallucinated product recommendations.
- Reduced LLM cost per turn by ~35% by restructuring prompts into a cacheable static prefix plus dynamic context injection and trimming redundant inter-agent payloads.
- Designed a new intent path end-to-end, including intent classification, structured NLU extraction schema (Pydantic), and planner routing, extending the assistant to a new class of user queries.
- Cut product-retrieval p95 latency from ~8s to ~4.5s, the largest single contributor to end-to-end discovery latency, as part of customer-launch readiness.
- Engineered real-time agentic chat over SSE/AG-UI/A2A with token-by-token streaming, cutting perceived response latency by 47%; implemented cross-channel context persistence so the planner retrieves prior conversation state when users switch surfaces (app to dotcom).
- Integrated CARE and AskBlue into a unified conversational journey with progressive login, increasing agentic-assisted engagement by ~20% and purchases by ~12%; built LLM-driven gifting recommendations that lifted conversion 15% during promotional periods.
- Led launch-readiness performance validation across planner, NLU, discovery, and presentation services at up to 32 RPS, diagnosing p95 and memory regressions from OpenTelemetry traces and log analysis, and eliminating a memory-creep trend across a 4-hour endurance test.
- Owned GCP deployment infrastructure for the Applied Machine Learning (AML) platform, configuring Cloud Run services and Cloud Deploy pipelines for agent rollouts across environments.
- Architected GCP data pipelines (BigQuery, Cloud Scheduler, Pub/Sub) processing 2M+ daily events at 99.9% reliability; built Java/Spring Boot services cutting API response times by 40%, and led blue-green/canary deployment practices maintaining 99.95% uptime under 3x Black Friday traffic.

Software Development Engineer

Nov 2023 – May 2025

Wissen Technology

- Clients: Goldman Sachs (Investment Banking), Spiro (Electric Vehicles)
- Engineered API integrations and optimized AWS infrastructure using Java Spring Boot, ensuring 99.9% uptime for post-trading settlement operations processing 8M+ daily trades.
- Built trade processing pipeline converting SII (A2V2) messages to XML instructions; implemented CI testing using FitNesse for automation and regression testing.
- Implemented multithreading with executor service and CompletableFutures, improving system performance by 70%; enhanced API security using JWT.
- Authored Jenkins CI/CD pipelines in Groovy, reducing deployment times by 80%; deployed Docker containers using Kubernetes on AWS EC2.
- Expanded EV platform customer base from 2 countries to 5, including Rwanda, Kenya, and Uganda, through backend system optimization.

Full Stack Engineer

Jul 2021 – Nov 2023

Tata Consultancy Services

- Optimized TCS iON Digital Marking Hub load times from 15s to 2s via asynchronous web APIs, serving 500+ organizations.
- Reduced API calls by 75% (10K to 2.5K monthly) via debouncing, improving search bar responsiveness by 50%.

- Implemented multi-language support (Hindi, Japanese, Arabic), increasing international user engagement by 30%.
- Achieved zero-defect release rate through CI/CD best practices and cross-functional collaboration using Jenkins.

PROJECTS

Ad Selection using Reinforcement Learning

- Developed RL model using Thompson Sampling to dynamically optimize ad selection, maximizing CTR through reward-based learning.

Placement Prediction Web Application

- Built ML application using Python and Logistic Regression achieving 89% accuracy; utilized by 10,000+ students with sub-second prediction times.

Hotel Review Sentiment Analysis

- Implemented NLP sentiment analysis pipeline with preprocessing optimizations, reducing dataset size by 20% and model execution time by 40%.

CERTIFICATIONS

- Machine Learning & Python Time Series (Udemy)
- Java & DSA (Scaler)
- HTML/CSS/JS (Microsoft)

AWARDS & HONORS

- Won Hackathon at the Advanced AI Engineering Campus Immersion (Scaler x CEC IIT Roorkee)
- Received iON Star of the Quarter Award for delivering project 2 weeks ahead of schedule with 15% efficiency increase
- Published research paper in IEEE on Impact of Covid-19 on Education Using Twitter Data at the 16th SMAP Conference, Greece

SKILLS

AI/ML: Generative AI, LLMs, Agent Orchestration, RAG, Google ADK, Vertex AI, NLP, Reinforcement Learning, NumPy

Cloud & Data: GCP, AWS, BigQuery, Pub/Sub, Cloud Run, Memcached

Observability: OpenTelemetry, Splunk, Grafana, Kibana, Dynatrace

Languages: Python, Java, JavaScript, SQL

Frameworks & Tools: Spring Boot, Microservices, REST API, Docker, Kubernetes, Jenkins

Databases: PostgreSQL, MySQL, MongoDB, Cassandra